

Geoscientific Machine Learning

Pankaj K Mishra

Table of Contents

1	Geoscientific Machine Learning	4
2	Getting started with Julia	6
2.1	Install Julia and VS Code	6
2.2	Three ways to run Julia code	7
2.3	Hello, Julia!	9
2.4	Your first calculations	9
2.5	Reusing code with functions	11
2.6	Making decisions and repeating things	12
2.7	Organising data with custom types	13
2.8	Getting comfortable with the REPL	13
2.9	Everyday patterns	14
2.10	Performance habits	15
2.11	Customising your Julia setup	17
2.12	The .julia directory: where packages live	18
2.13	Using Julia behind a proxy	19
2.14	Reproducible environments	21
3	Files & Data Manipulation	26
3.1	Writing and reading text files	26
3.2	Delimited data with DelimitedFiles	28
3.3	CSV files and DataFrames	30
3.4	Working with file paths	35
3.5	Searching, copying, and renaming files	36
3.6	Geoscientific file formats	41
3.7	Clean up	44
4	Plotting & Data Visualisation	45
4.1	The Makie ecosystem	45
4.2	Installing CairoMakie	46
4.3	Your first plot	47
4.4	Scatter plots	49
4.5	Line plots with multiple series	51
4.6	Error bars	52
4.7	Bar charts and histograms	53

4.8	Heatmaps and contour plots	55
4.9	Multi-panel figures (subplots)	58
4.10	Customising appearance	60
4.11	Log-scale axes	63
4.12	Annotations and markers	64
4.13	Saving figures	65
4.14	Putting it together: a geoscience figure	66
4.15	Geographic maps with GeoMakie	67
4.16	Where to learn more	70
5	Neural Networks	71
5.1	Your first neural network: Sauna Satisfaction Predictor	71
5.2	What is a neural network?	77

1 Geoscientific Machine Learning

Welcome to **Geoscientific Machine Learning**.

If you're reading this, you probably already have some familiarity with geoscience. You're used to thinking in terms of physical processes, models, scales, and uncertainty. When it comes to computation, though, you may have mostly worked through existing tools: running models, adjusting parameters, or modifying scripts written by someone else. That's a common place to be, and it works, until you want to change the model itself or test an idea that doesn't fit into the existing workflow.

You might also be here because you keep hearing about AI and machine learning and you're not sure what to make of it. Is it actually useful for your research, or is it just hype? Will it help you work faster, or will it send you down a rabbit hole? Can it help you do things you thought were out of reach before, or is it mostly a waste of time?

Before you decide, it's worth trying it in a way that respects how geoscience works. That's what this book is about.

By **scientific machine learning**, we mean instead of training a model only on data and hoping it behaves, you guide the learning with what you already know: physical laws, conservation principles, symmetries, and numerical structure. The goal is not just to fit observations, but to produce models that behave sensibly when you change conditions or run them for longer times. By **geoscientific machine learning**, we mean applying scientific machine learning to geoscience problems. That means combining geoscience domain knowledge with learning methods so the models respect physics, scale relationships, and observational constraints.

If you don't yet have a favorite programming language, or you're still deciding what's worth learning, here's the important part: the exact language matters less than having *one*. You need a way to talk to computers clearly, so they can help you turn ideas into experiments and hypotheses into something you can test.

In this book we'll choose **Julia**. If you spend time with it, you'll see why—especially once you start doing more demanding computation. We'll begin slowly and keep things practical, using Julia to express ideas you already know. Machine learning comes later where it connects naturally geoscience practices.

If you're new to programming, expect some friction at first (we'll try to minimise that). That's normal. The payoff is that you gain more control over your models and more confidence in the results they produce.

i This book is under construction

Some sections are incomplete. If there's something you'd like to see included, open an issue and let me know what you're working on.

2 Getting started with Julia

You are starting from scratch, and that is perfectly fine. No prior programming experience is assumed.

In this chapter we will install Julia and VS Code, write our first lines of code, and build up to functions, loops, and custom types. By the end you will be comfortable running Julia interactively, saving scripts, and managing packages in reproducible environments. If you have used another language before (Python, R, MATLAB), you will notice how little setup Julia needs to get going.

2.1 Install Julia and VS Code

You need two things: **Julia** (the language) and **VS Code** (the editor you will write code in). Install them in this order:

1. Install **VS Code** from code.visualstudio.com.
2. Install **Julia** using **juliaup**, a small tool that manages Julia versions for you:

- **Windows (PowerShell)**

```
winget install julia -s msstore  
juliaup add release
```

- **macOS (Terminal)**

```
curl -fsSL https://install.julialang.org | sh  
juliaup add release
```

- **Linux (Terminal)**

```
curl -fsSL https://install.julialang.org | sh  
juliaup add release
```

3. Open a new terminal and check that Julia is available:

```
julia --version
```

If it prints something like `julia version 1.11.2`, you are ready.

4. Open VS Code, go to the **Extensions** panel (the square icon on the left sidebar, or press `Ctrl/Cmd+Shift+X`), and search for **Julia**. Install the one published by **julia**lang. This single extension gives you syntax highlighting, code execution, an integrated REPL, a debugger, and a plot viewer. That is everything you need to get started.

Reading code blocks in this book

Throughout this book you will see two styles of code block. **Julia code** has a thick coloured bar on the left edge:

```
println("I am Julia code")
```

Shell commands, configuration files, and other non-Julia text have a plain border with no coloured bar:

```
echo "I am a shell command"
```

If a block has the coloured side bar, you can type it into the Julia REPL or a `.jl` file. If it does not, it is meant for your terminal, a config file, or is just illustrative output.

2.2 Three ways to run Julia code

Before we start writing code, it helps to know that there are three ways to run Julia. Each is useful in different situations:

2.2.1 1. Inside VS Code (recommended for beginners)

This is the easiest way. Create a file ending in `.jl` (e.g., `test.jl`), type your code, and press **Ctrl+Enter** (Windows/Linux) or **Cmd+Enter** (macOS) to run one line or a selected block. The Julia extension opens a built-in REPL panel at the bottom of VS Code and sends your code there. You see the results immediately, and you can go back and edit your file at any time.

This is what we will use throughout this book. It combines the best of both worlds: your code is saved in a file you can come back to, but you can run pieces of it interactively.

2.2.2 2. Interactively in the REPL (command line)

Open a terminal and type `julia` to start the **REPL** (Read-Eval-Print Loop). You get a `julia>` prompt where you can type expressions and see results instantly:

```
$ julia
```

```

      _
     _ _(_)- |
    (-)      | (-) (-) | Version 1.11.2
   _ _ _ | | _ _ _ _ |
  | | | | | | | / - ` | |
  | | | - | | | | (- | | |
 - / | \ -- ' - | - | \ -- ' - |
 | -- / |
```

```
julia> 2 + 3
5
```

```
julia> sqrt(144)
12.0
```

The REPL is great for quick experiments: testing a formula, checking how a function works, or exploring a dataset. But nothing is saved to a file, so it is not ideal for work you want to keep.

2.2.3 3. Running a script from the terminal

If you have a file called `myscript.jl`, you can run the whole thing at once from the terminal:

```
julia myscript.jl
```

Julia reads the file from top to bottom, executes every line, prints any output, and exits. This is how you run production code, batch jobs on a cluster, or automated workflows. You don't see intermediate results, only what the script explicitly prints.

2.2.4 Which one should I use?

Situation	Best option
Learning and exploring	VS Code (Ctrl/Cmd+Enter) or REPL

Situation	Best option
Developing code you want to keep	VS Code (code is in a file, but you run it interactively)
Quick one-off calculations	REPL
Running a finished analysis on a cluster	<code>julia myscript.jl</code>
Automated or scheduled jobs	<code>julia myscript.jl</code>

You will naturally switch between these as you get more comfortable. For now, open VS Code and follow along.

2.3 Hello, Julia!

Let's make sure everything works. In VS Code, create a new file called `test.jl` and type the following:

```
println("Hello, Julia!")
x = [1, 2, 3, 4, 5]
println("Sum: ", sum(x))
```

Press **Ctrl+Enter** (Windows/Linux) or **Cmd+Enter** (macOS) to run each line. You should see:

```
Hello, Julia!
Sum: 15
```

Congratulations, you just ran your first Julia program. Everything after this point builds on what you just did.

2.4 Your first calculations

Julia works like a calculator right out of the box. You don't need to import anything or set anything up:

```
1 + 2, sin(1.0)
```

```
(3, 0.8414709848078965)
```

This computed $1 + 2$ and the sine of 1.0 (in radians) and returned both results as a pair. The comma groups them into a **tuple**, just a way of showing multiple answers together.

2.4.1 Storing values in variables

In programming, a **variable** is a name you give to a value so you can use it later. Think of it as a sticky note on a number:

```
a = 10
b = 3
a + b, a - b, a * b, a / b
```

(13, 7, 30, 3.3333333333333335)

Here *a* is a label for 10 and *b* is a label for 3. The next line asks Julia to add, subtract, multiply, and divide them, all at once.

You can name variables almost anything: *temperature*, *depth_km*, *n_samples*. Use names that describe what the value means. Your future self will thank you.

2.4.2 Collecting values in a list

In real work you almost never deal with just one number. You have a list of measurements, a column of data, a time series. In Julia, a list of values is called an **Array** (or **Vector** when it is one-dimensional). You create one with square brackets:

```
x = [1, 2, 3, 4, 5]
sum(x), length(x)
```

(15, 5)

`sum(x)` adds up all the values. `length(x)` counts how many there are. No imports, no setup.

2.4.3 What about text?

Not everything is a number. Julia handles text (called **strings**) too. Strings are wrapped in double quotes:

```
name = "Julia"
println("Hello, $name!")
```

The `$` inside a string inserts the value of a variable. This is called **string interpolation**. You can also insert expressions: `"2 + 3 = $(2 + 3)"`.

2.5 Reusing code with functions

If you find yourself doing the same calculation over and over, wrap it in a **function**. A function is like a recipe: you give it inputs (ingredients) and it gives you back a result (the dish).

The simplest way to define a function is the one-liner:

```
f(x) = x^2 + 2x + 1
f(3)
```

16

This says: “define `f` so that it takes `x`, squares it, adds twice `x`, adds 1, and returns the result.” Now `f(3)` gives 16, `f(10)` gives 121, and so on.

For anything longer than one line, use the multi-line form:

```
function kinetic_energy(mass, velocity)
    return 0.5 * mass * velocity^2
end

kinetic_energy(70.0, 3.0)  # a 70 kg person walking at 3 m/s
```

Functions can also return more than one value. Unpack them into separate variables:

```
function minmax(x)
    return minimum(x), maximum(x)
end

lo, hi = minmax([4, 1, 7, 2])
println("min = $lo, max = $hi")
```

 Why functions matter

Julia's compiler makes code inside functions run very fast, often as fast as C or Fortran. Code typed at the top level (outside a function) cannot be optimised the same way. So the habit of wrapping your work in functions is not just about tidiness, it is about speed.

2.6 Making decisions and repeating things

2.6.1 If / else: choosing what to do

Programs often need to choose what to do based on a condition. In Julia, this reads almost like English:

```
x = 5
if x > 0
    println("positive")
elseif x < 0
    println("negative")
else
    println("zero")
end
```

positive

`println` means “print this line of text.” Julia checks the conditions from top to bottom and runs the first one that is true, then skips the rest.

2.6.2 Loops: repeating work

A **loop** lets you repeat something many times without copying and pasting. The most common loop says “for each value in this range, do something”:

```
for i in 1:5
    println(i^2)
end
```

```
1
4
9
16
25
```

`1:5` means the numbers 1, 2, 3, 4, 5. For each one, Julia squares it and prints the result. In geoscience, you might use a loop to process every station in a survey, every layer in a model, or every time step in a simulation.

2.7 Organising data with custom types

A geoscience measurement is rarely just a single number. It has coordinates, a value, an uncertainty, maybe a timestamp. Julia lets you bundle related values together using a **struct** (short for structure):

```
struct Point
    x::Float64
    y::Float64
end

p = Point(1.0, 2.0)
p.x, p.y
```

```
(1.0, 2.0)
```

This creates a new type called `Point` with two fields. `Float64` means “a decimal number with double precision.” Once defined, you create a point by writing `Point(1.0, 2.0)` and access its parts with `p.x` and `p.y`.

Why bother with structs?

As your projects grow, structs keep things organised. Instead of passing around loose variables like `lat`, `lon`, `elevation`, you pass a single `Station` object that holds all of them together. Julia’s compiler also uses the type information to make your code faster.

2.8 Getting comfortable with the REPL

The **REPL** (Read-Eval-Print Loop) is the interactive prompt you see when you type `julia` in a terminal. It is a great place to try things out quickly.

2.8.1 Built-in modes

The REPL has several modes. You switch by pressing a single key at the start of an empty line:

Key	Mode	What it does
<code>]</code>	Pkg	Package manager: add, remove, update packages
<code>;</code>	Shell	Run shell commands without leaving Julia
<code>?</code>	Help	Look up documentation for any function or type
Backspace	Julia	Return to normal Julia mode

Try it: press `?` then type `sum` and hit Enter. Julia shows you the documentation for `sum`, with examples.

2.8.2 Tab completion and Unicode

Press **Tab** to autocomplete names. This works for functions, variables, file paths, and even Unicode:

- Type `pri` then **Tab** → `print`
- Type `\alpha` then **Tab** → α
- Type `\sqrt` then **Tab** → $\sqrt$

This is how you type mathematical symbols in Julia code. Any LaTeX name works.

2.8.3 Handy shortcuts

- **Semicolons suppress output:** `x = rand(1000);` assigns the array without printing all 1000 numbers.
- **ans holds the last result:** after typing `2 + 3`, typing `ans * 10` gives 50.
- **Underscores in numbers:** write `1_000_000` instead of `1000000` for readability. Julia ignores the underscores.

2.9 Everyday patterns

These are practical patterns you will use constantly.

2.9.1 Broadcasting with the dot `.`

Apply any function or operator to every element of an array by adding a dot:

```
x = [1, 2, 3, 4]
sin.(x)           # sine of each element
x .^ 2           # square each element
x .+ 10          # add 10 to each element
```

This works with your own functions too:

```
f(x) = x^2 + 1
f.([1, 2, 3])    # returns [2, 5, 10]
```

2.9.2 Finding documentation and methods

Use `methods()` to see all versions of a function, and `@which` to find out which version gets called:

```
methods(+)        # all methods for +
@which 1.0 + 2.0   # which method handles Float64 + Float64
```

2.9.3 Timing your code

Use `@time` for a quick measurement. Run it **twice** because the first call includes compilation time:

```
@time sum(rand(1_000_000))  # first call: includes compilation
@time sum(rand(1_000_000))  # second call: actual speed
```

For serious benchmarking, use the `BenchmarkTools` package:

```
using BenchmarkTools
@btime sum(rand(1_000_000))
```

2.10 Performance habits

These two habits will save you hours of debugging slow code.

2.10.1 Put work inside functions

Code at the top level (the global scope) runs slowly because Julia cannot predict the types of global variables. The same code inside a function runs orders of magnitude faster:

```
# Slow - top level
x = rand(1_000_000)
total = 0.0
for val in x
    total += val
end

# Fast - inside a function
function mysum(x)
    total = 0.0
    for val in x
        total += val
    end
    return total
end
mysum(rand(1_000_000))
```

2.10.2 Loop in the right order

Julia stores matrices **column by column** in memory (like Fortran), not row by row (like C or Python). When looping over a matrix, always loop over rows in the inner loop:

```
A = rand(1000, 1000)

# Fast - columns in the outer loop, rows in the inner loop
for j in 1:1000      # column
    for i in 1:1000  # row
        A[i, j] += 1
    end
end
```


2.11 Customising your Julia setup

2.11.1 Startup file

If you load the same packages every session, add them to your startup file. Julia runs this file automatically on launch.

The file lives at `~/.julia/config/startup.jl`. Create the directory and file if they don't exist:

```
mkdir -p ~/.julia/config
echo 'try using Revise catch end' > ~/.julia/config/startup.jl
```

On Windows (PowerShell):

```
New-Item -ItemType Directory -Force -Path "$env:USERPROFILE\.julia\config"
Set-Content "$env:USERPROFILE\.julia\config\startup.jl" 'try using Revise catch end'
```

Revise.jl

Revise watches your source files and automatically picks up changes without restarting Julia. Install it once with `] add Revise`, then load it from your startup file as shown above. It is one of the most useful Julia packages.

2.11.2 Point VS Code to Julia manually

In most cases VS Code finds Julia automatically. If it does not, go to **Settings** → **Julia: Executable Path** and enter the path:

- **Windows**

`%LOCALAPPDATA%\Programs\Julia-*\bin\julia.exe`

or

`%USERPROFILE%\.juliaup\bin\julia.exe`

- **macOS**

`/Applications/Julia-*.app/Contents/Resources/julia/bin/julia`

or

```
$HOME/.juliaup/bin/julia
```

- **Linux**

```
/usr/bin/julia
```

or

```
$HOME/.juliaup/bin/julia
```

2.12 The `.julia` directory: where packages live

When you install Julia and start adding packages, Julia creates a hidden directory called `.julia` in your home folder:

OS	Default location
Windows	<code>C:\Users\YourName\.julia\</code>
macOS	<code>/Users/YourName/.julia/</code>
Linux	<code>/home/YourName/.julia/</code>

This directory is called the **depot**. It stores everything Julia needs:

Subfolder	What's inside
<code>packages/</code>	Source code of every package you install
<code>compiled/</code>	Precompiled versions of those packages
<code>artifacts/</code>	Binary libraries and datasets that packages download
<code>registries/</code>	The General registry, Julia's index of all public packages
<code>logs/</code>	Build and install logs
<code>environments/</code>	Project environments and their <code>Manifest.toml</code> files
<code>config/</code>	Your <code>startup.jl</code> file

2.12.1 Why it grows large

Every package brings its own dependencies, compiled caches, and sometimes large binary artefacts. A plotting library like CairoMakie can pull in hundreds of megabytes. Over time, `.julia` can easily reach 5–20 GB.

On a personal laptop this is fine. But on **university HPC clusters and shared Linux servers**, your home directory often has a quota, sometimes as little as 5 or 10 GB. If `.julia` fills it up, you will see “disk quota exceeded” errors and Julia will stop installing packages.

2.12.2 Moving .julia to a larger disk

Set the JULIA_DEPOT_PATH environment variable to point somewhere with more space:

- **Windows (PowerShell, persist permanently)**

```
setx JULIA_DEPOT_PATH "D:\JuliaDepot"
```

- **macOS / Linux** — add to ~/.bashrc or ~/.zshrc:

```
export JULIA_DEPOT_PATH="$HOME/julia-depot"
```

- **On a cluster** (use scratch or project space):

```
export JULIA_DEPOT_PATH="/scratch/$USER/julia-depot"
```

💡 Already have a large .julia in your home?

Move the existing directory first, then set the variable:

```
mv ~/.julia /scratch/$USER/julia-depot  
export JULIA_DEPOT_PATH="/scratch/$USER/julia-depot"
```

On Windows (PowerShell):

```
Move-Item "$env:USERPROFILE\.julia" "D:\JuliaDepot"  
setx JULIA_DEPOT_PATH "D:\JuliaDepot"
```

2.13 Using Julia behind a proxy

On university and corporate networks, internet traffic often goes through a proxy. Julia needs to know about this to download packages. If you skip this step, `Pkg.add(...)` will hang or time out.

Ask your IT department for the proxy address. Then set it as an environment variable:

Windows (PowerShell, persist permanently)

```
setx HTTP_PROXY "http://proxy.example.com:8080"  
setx HTTPS_PROXY "http://proxy.example.com:8080"
```

Close and reopen your terminal after running setx.

macOS / Linux — add to ~/.bashrc or ~/.zshrc:

```
export HTTP_PROXY="http://proxy.example.com:8080"  
export HTTPS_PROXY="http://proxy.example.com:8080"
```

Then reload:

```
source ~/.zshrc # or ~/.bashrc
```

If authentication is required, include credentials in the URL:

```
http://username:password@proxy.example.com:8080
```

2.13.1 Verify it works

Open the Julia REPL and try:

```
ENV["HTTP_PROXY"]      # should print your proxy address  
using Pkg  
Pkg.add("Example")      # should download successfully
```

💡 No-proxy for local addresses

If internal servers should bypass the proxy:

```
export NO_PROXY="localhost,127.0.0.1,.internal.example.com"
```

On Windows (PowerShell):

```
$env:NO_PROXY = "localhost,127.0.0.1,.internal.example.com"
```

2.14 Reproducible environments

Imagine you wrote a script that reads geoscientific data, runs an inversion, and produces a beautiful map. It works perfectly today. Six months later your colleague tries to run it and gets errors everywhere. A package was updated, a function was renamed, a dependency now needs a newer Julia version. This is the “*it works on my machine*” problem, and it is one of the biggest obstacles to reproducible science.

Julia’s **built-in package manager** solves this at the language level. Every Julia project can carry two small text files that record *exactly* which packages (and which versions) were used. Anyone, anywhere, can recreate the identical environment with a single command.

2.14.1 What a Julia project looks like

A typical Julia project is just a folder with a few files:

```
MyProject/
├── Project.toml      # what you depend on
├── Manifest.toml     # exact snapshot of every version
└── src/
    └── MyProject.jl  # your code
```

You may also have a `test/`, `docs/`, or `scripts/` folder, but the two `.toml` files are the heart of reproducibility.

2.14.2 Project.toml: what your project needs

This file lists the **direct dependencies**, the packages *you* asked for. Here is the `Project.toml` for this very book:

```
[deps]
CSV = "336ed68f-0bac-5ca0-87d4-7b16caf5d00b"
CairoMakie = "13f3f980-e62b-5c42-98c6-ff1f3baf88f0"
DataFrames = "a93c6f00-e57d-5684-b7b6-d8193f3e46c0"
DelimitedFiles = "8bb1440f-4735-579b-a4ab-409b98df4dab"
Statistics = "10745b16-79ce-11e8-11f9-7d13ad32a3b2"
```

Each entry has a name and a **UUID** (a unique identifier so Julia never confuses two packages with the same name). You rarely type UUIDs by hand; the package manager fills them in when you run `Pkg.add("CSV")`.

You can optionally add a `[compat]` section to pin version ranges:

```
[compat]
CSV = "0.10"
DataFrames = "1.6"
julia = "1.10"
```

This tells Julia: “My code is tested with CSV 0.10.x, DataFrames 1.6.x, and Julia 1.10 or newer.” It prevents someone from accidentally installing an incompatible future version.

2.14.3 Manifest.toml: the exact snapshot

When you install packages, Julia resolves every dependency (and every dependency’s dependency) and writes the result into `Manifest.toml`. This file records the **exact version** of every single package in the tree. It can be hundreds of lines long, and that is normal.

Think of it this way:

File	Analogy	You edit it?
<code>Project.toml</code>	A recipe’s ingredient list (“flour, eggs, sugar”)	Yes
<code>Manifest.toml</code>	The supermarket receipt with exact brands and batch numbers	No (generated automatically)

💡 Should you commit `Manifest.toml` to Git?

For applications and papers: yes. It guarantees bit-for-bit reproducibility. Anyone who checks out your repo and runs `Pkg.instantiate()` gets exactly the same package versions.

For libraries (reusable packages): usually no. You want downstream users to resolve versions against *their* environment, not yours.

2.14.4 Creating and using environments

Start a new project in any empty folder:

```
using Pkg
Pkg.activate(".") # use the current folder as the environment
Pkg.add("CSV")    # adds CSV to Project.toml and resolves Manifest.toml
```

```
Pkg.add("DataFrames")
```

Or in the REPL's Pkg mode (press]):

```
(@v1.11) pkg> activate .  
(MyProject) pkg> add CSV DataFrames
```

Julia creates `Project.toml` and `Manifest.toml` for you.

Reproduce someone else's environment:

```
using Pkg  
Pkg.activate(".")          # point to the folder containing their Project.toml  
Pkg.instantiate()          # download the exact versions from Manifest.toml
```

That is it. Two commands and you have an identical setup. No guessing, no version conflicts.

2.14.5 The default (global) environment

When you start Julia without activating a project, you are in the **global environment** (shown as `@v1.11` in the REPL prompt). Packages you add here are available everywhere, which is convenient for everyday tools like `Revise` or `BenchmarkTools`. But for actual research work, always create a **project-specific environment** so your results are reproducible.

How other languages handle this

If you come from **Python**, Julia's `Project.toml` is similar to `requirements.txt` or `pyproject.toml`, and `Manifest.toml` is like a `pip freeze` snapshot or a Poetry/PDM lock file. The difference is that Julia's package manager is built into the language, with no separate tool to install.

If you come from **C++**, reproducible builds are traditionally much harder. You might use `CMake` with pinned versions, `Conan`, or `vcpkg`, but there is no single built-in standard. Julia's approach is closer to Rust's `Cargo.toml` / `Cargo.lock` system.

In all cases the idea is the same: **separate what you want from what you got**, so someone else can recreate the exact same setup later.

2.14.6 Key Pkg commands explained

Here is what the most common Pkg functions actually do:

Pkg.activate(".") Switch into a project environment. This tells Julia: “use the `Project.toml` in this folder instead of the global one.” The REPL prompt changes from `(@v1.11)` to the project name. Nothing is installed yet; you are just pointing Julia at a folder.

Pkg.add("CSV") Add a package. Writes the package name and UUID into `Project.toml`, resolves a compatible set of versions, downloads everything needed, and records the result in `Manifest.toml`. If `Project.toml` does not exist yet, Julia creates it for you.

Pkg.instantiate() Recreate an environment from existing files. If a `Manifest.toml` is present, it downloads the *exact* versions listed there. If only a `Project.toml` exists (no manifest), it resolves fresh versions and creates the manifest. This is the command you run after cloning someone else’s repository.

Pkg.resolve() Re-resolve versions without adding or removing anything. Useful after you edit `Project.toml` by hand (for example, adding a `[compat]` entry). It updates `Manifest.toml` to match the current `Project.toml` without downloading new packages.

Pkg.update() Update all packages to the newest versions that are still compatible with your `[compat]` constraints. This rewrites `Manifest.toml`. Run this periodically to pick up bug fixes and performance improvements.

Pkg.status() List installed packages and their versions. In Pkg-mode just type `st` or `status`.

Pkg.pin("CSV") Lock a package at its current version so that `Pkg.update()` will not change it. Useful when you know a newer release breaks your workflow. Unpin with `Pkg.free("CSV")`.

Pkg.rm("CSV") Remove a package from the project. Deletes it from `Project.toml` and re-resolves the manifest.

Pkg.gc() Garbage-collect. Removes cached packages that are no longer used by any environment on your machine. Helpful when your `.julia` folder grows large (see the earlier section on the `.julia` directory).

Pkg.build() Re-run build scripts for installed packages. Occasionally needed after system-level changes (new compiler, updated shared libraries).

2.14.7 Quick reference

Task	REPL Pkg-mode (<code>]</code>)	Script
Activate an environment	<code>activate .</code>	<code>Pkg.activate(".")</code>
Add a package	<code>add CSV</code>	<code>Pkg.add("CSV")</code>
Remove a package	<code>rm CSV</code>	<code>Pkg.rm("CSV")</code>
Recreate from lock file	<code>instantiate</code>	<code>Pkg.instantiate()</code>
Re-resolve after hand edits	<code>resolve</code>	<code>Pkg.resolve()</code>
See what is installed	<code>status</code>	<code>Pkg.status()</code>
Update all packages	<code>update</code>	<code>Pkg.update()</code>

Task	REPL Pkg-mode (<code>()</code>)	Script
Pin a version	<code>pin CSV@0.10.14</code>	<code>Pkg.pin("CSV", v"0.10.14")</code>
Unpin	<code>free CSV</code>	<code>Pkg.free("CSV")</code>
Clean unused caches	<code>gc</code>	<code>Pkg.gc()</code>
Re-run build scripts	<code>build</code>	<code>Pkg.build()</code>

 Learn more

The official **Pkg documentation** covers everything in detail, including custom registries, private packages, artifacts, and more:

<https://pkgdocs.julialang.org/>

With these two files and a handful of commands, your work stays reproducible, whether you revisit it next year or share it with a colleague on the other side of the world.

3 Files & Data Manipulation

At this point you know how to install Julia, run code in VS Code or the REPL, write functions, and create reproducible environments with `Project.toml`.

In this chapter we will put those skills to work on something you will do constantly in geoscience: **reading data from files, reshaping it, and writing results back out**. We start with plain text files, move on to delimited numeric data, then to CSV tables with the DataFrames ecosystem, and finish with a tour of specialised geoscientific formats (NetCDF, HDF5, SEG-Y, Shapefiles). We create small example files right here in the code so you can run every block and see the results yourself.

If some of the syntax looks unfamiliar at first, especially the `do ... end` pattern or the $\Rightarrow$ arrow, don't worry. These are common Julia patterns that will become second nature with practice. The important thing is to understand *what* each block does, not to memorise the syntax right now.

3.1 Writing and reading text files

The simplest kind of file is plain text. Let's create one:

```
# Create a small text file
open("stations.txt", "w") do f
    write(f, "Station  Latitude  Longitude  Elevation_m\n")
    write(f, "AAL      64.35    25.75    120\n")
    write(f, "BBR      61.50    23.80    95\n")
    write(f, "CCK      60.17    24.94    15\n")
end
```

33

What happened here? `open("stations.txt", "w")` creates (or overwrites) a file named `stations.txt`. The `"w"` means “write mode.” The `do f ... end` block is Julia's way of saying “while the file is open, call it `f` and do these things, then close it automatically.” Each `write` call puts a line of text into the file.

Now let's read it back:

```
content = read("stations.txt", String)
println(content)
```

Station	Latitude	Longitude	Elevation_m
AAL	64.35	25.75	120
BBR	61.50	23.80	95
CCK	60.17	24.94	15

`read("stations.txt", String)` loads the entire file into a single string. For a small file like this, that is perfectly fine.

3.1.1 Reading line by line

Sometimes you want to process a file one line at a time, for instance to skip a header or to parse each row:

```
open("stations.txt", "r") do f
  for line in eachline(f)
    println("→ ", line)
  end
end
```

```
→ Station Latitude Longitude Elevation_m
→ AAL      64.35      25.75      120
→ BBR      61.50      23.80      95
→ CCK      60.17      24.94      15
```

The "r" means "read mode." `eachline(f)` gives you one line at a time, which is memory-efficient even for very large files.

3.1.2 Splitting lines into columns

In geoscience data files, columns are usually separated by spaces, tabs, or commas. You can split a line into pieces with `split()`:

```
open("stations.txt", "r") do f
  header = readline(f)           # read and skip the header
  for line in eachline(f)
```

```

    parts = split(line)          # splits on whitespace by default
    name  = parts[1]
    lat   = parse(Float64, parts[2])
    lon   = parse(Float64, parts[3])
    elev  = parse(Float64, parts[4])
    println("$name is at ($lat, $lon), elevation $elev m")
end
end

```

```

AAL is at (64.35, 25.75), elevation 120.0 m
BBR is at (61.5, 23.8), elevation 95.0 m
CCK is at (60.17, 24.94), elevation 15.0 m

```

`split(line)` breaks the string into a list of substrings. `parse(Float64, ...)` converts a text string like "64.35" into an actual number. This pattern (read a line, split it, parse the pieces) is the bread and butter of working with geoscience text files.

3.2 Delimited data with `DelimitedFiles`

For files with a regular grid of numbers (like gravity measurements or temperature readings), Julia has a built-in module called `DelimitedFiles` that reads the whole file into a matrix in one step.

Let's create a small data file and read it:

```

# Create a space-delimited data file (no header, just numbers)
open("gravity.txt", "w") do f
    write(f, "64.35 25.75 -12.3\n")
    write(f, "64.40 25.80 -11.8\n")
    write(f, "64.45 25.85 -13.1\n")
    write(f, "64.50 25.90 -12.7\n")
    write(f, "64.55 25.95 -14.0\n")
end

```

20

```

using DelimitedFiles

data = readdlm("gravity.txt")    # reads into a matrix

```

```
println("Size: ", size(data))      # 5 rows, 3 columns
```

Size: (5, 3)

Now you can pull out individual columns:

```
lat = data[:, 1]      # all rows, first column
lon = data[:, 2]      # all rows, second column
gz  = data[:, 3]      # all rows, third column

println("Latitudes: ", lat)
println("Gravity anomalies: ", gz)
```

Latitudes: [64.35, 64.4, 64.45, 64.5, 64.55]

Gravity anomalies: [-12.3, -11.8, -13.1, -12.7, -14.0]

The `[:, 1]` notation means “all rows, column 1.” This is the same notation used in MATLAB and NumPy.

You can also write a matrix back out:

```
# Add 1.0 to all gravity values and save
gz_corrected = gz .+ 1.0
output = hcat(lat, lon, gz_corrected)  # glue columns side by side
writedlm("gravity_corrected.txt", output, '\t')  # tab-separated

# Verify
println(read("gravity_corrected.txt", String))
```

```
64.35  25.75  -11.3
64.4    25.8   -10.8
64.45   25.85  -12.1
64.5    25.9   -11.7
64.55   25.95  -13.0
```

`hcat` stands for “horizontal concatenate”. It sticks columns together into a matrix. `writedlm` writes the matrix to a file, using the tab character `'\t'` as the separator.

3.3 CSV files and DataFrames

CSV (Comma-Separated Values) is probably the most common data format in science. Julia has excellent CSV support through the `CSV.jl` and `DataFrames.jl` packages.

Installing packages

If you haven't installed these packages yet, press `]` in the REPL to enter Pkg mode and type:

```
add CSV DataFrames
```

Or from normal Julia code: `using Pkg; Pkg.add(["CSV", "DataFrames"])`. You only need to do this once.

3.3.1 Creating a DataFrame

A **DataFrame** is a table: rows and columns, like a spreadsheet. Each column has a name and all values in a column have the same type. If you have used pandas in Python, R's `data.frame`, or even Excel, the concept is the same.

```
using DataFrames

# Create a table of borehole samples
samples = DataFrame(
    borehole = ["BH-01", "BH-01", "BH-01", "BH-02", "BH-02"],
    depth_m  = [10.0, 20.0, 30.0, 15.0, 25.0],
    rock_type = ["granite", "granite", "gneiss", "granite", "schist"],
    density  = [2.65, 2.68, 2.75, 2.63, 2.71]
)
```

	borehole	depth_m	rock_type	density
	String	Float64	String	Float64
1	BH-01	10.0	granite	2.65
2	BH-01	20.0	granite	2.68
3	BH-01	30.0	gneiss	2.75
4	BH-02	15.0	granite	2.63
5	BH-02	25.0	schist	2.71

Julia displays the table neatly. Each column is a named array, and you can work with them individually or together.

3.3.2 Writing and reading CSV

Let's save this table to a CSV file and read it back:

```
using CSV

# Write to CSV
CSV.write("samples.csv", samples)

# Read it back
samples_loaded = CSV.read("samples.csv", DataFrame)
```

	borehole	depth_m	rock_type	density
	String7	Float64	String7	Float64
1	BH-01	10.0	granite	2.65
2	BH-01	20.0	granite	2.68
3	BH-01	30.0	gneiss	2.75
4	BH-02	15.0	granite	2.63
5	BH-02	25.0	schist	2.71

That's it. One line to write, one line to read. CSV.jl automatically detects column types, handles headers, and deals with missing values.

3.3.3 Selecting columns

Grab one or more columns by name:

```
# One column - returns a vector
samples.depth_m
```

5-element Vector{Float64}:

```
10.0
20.0
30.0
15.0
25.0
```

```
# Multiple columns - returns a new DataFrame
select(samples, :borehole, :density)
```

	borehole	density
	String	Float64
1	BH-01	2.65
2	BH-01	2.68
3	BH-01	2.75
4	BH-02	2.63
5	BH-02	2.71

The `:` before a column name makes it a **Symbol**, Julia's way of referring to names. You will see this pattern everywhere in DataFrames.

3.3.4 Filtering rows

Keep only rows that match a condition:

```
# Samples deeper than 20 m
filter(row → row.depth_m > 20, samples)
```

	borehole	depth_m	rock_type	density
	String	Float64	String	Float64
1	BH-01	30.0	gneiss	2.75
2	BH-02	25.0	schist	2.71

The `row → row.depth_m > 20` part is an **anonymous function**, a small throwaway function that takes one argument (`row`) and returns true or false. Don't worry about the syntax too much; think of it as saying "keep rows where depth is greater than 20."

```
# Only granite samples
filter(row → row.rock_type == "granite", samples)
```

	borehole	depth_m	rock_type	density
	String	Float64	String	Float64
1	BH-01	10.0	granite	2.65
2	BH-01	20.0	granite	2.68
3	BH-02	15.0	granite	2.63

3.3.5 Adding and transforming columns

```
# Add a new column
samples.weight_kg = samples.density .* 1000.0 # density × 1000
samples
```

	borehole	depth_m	rock_type	density	weight_kg
	String	Float64	String	Float64	Float64
1	BH-01	10.0	granite	2.65	2650.0
2	BH-01	20.0	granite	2.68	2680.0
3	BH-01	30.0	gneiss	2.75	2750.0
4	BH-02	15.0	granite	2.63	2630.0
5	BH-02	25.0	schist	2.71	2710.0

The `.*` (dot-star) applies the multiplication to every row. You saw this broadcasting pattern in Chapter 1.

3.3.6 Grouping and summarising

This is where DataFrames really shine. Group your data by a category and compute summaries:

```
using Statistics

# Average density per borehole
gdf = groupby(samples, :borehole)
combine(gdf, :density => mean => :avg_density)
```

	borehole	avg_density
	String	Float64
1	BH-01	2.69333
2	BH-02	2.67

Read this as: “group the table by `:borehole`, then for the `:density` column compute the mean and call the result `:avg_density`.” The `=>` arrows connect input $\rightarrow$ function $\rightarrow$ output name.

```
# Multiple summaries at once
combine(
    groupby(samples, :rock_type),
    :density => mean => :avg_density,
    :density => minimum => :min_density,
```

```
:depth_m => maximum => :deepest
)
```

	rock_type	avg_density	min_density	deepest
	String	Float64	Float64	Float64
1	granite	2.65333	2.63	20.0
2	gneiss	2.75	2.75	30.0
3	schist	2.71	2.71	25.0

3.3.7 Joining tables

In real projects your data is often spread across multiple files. Joining brings them together:

```
# A second table with borehole coordinates
locations = DataFrame(
    borehole = ["BH-01", "BH-02", "BH-03"],
    latitude = [64.35, 64.40, 64.45],
    longitude = [25.75, 25.80, 25.85]
)

# Inner join - only boreholes that appear in both tables
innerjoin(samples, locations, on = :borehole)
```

	borehole	depth_m	rock_type	density	weight_kg	latitude	longitude
	String	Float64	String	Float64	Float64	Float64	Float64
1	BH-01	10.0	granite	2.65	2650.0	64.35	25.75
2	BH-01	20.0	granite	2.68	2680.0	64.35	25.75
3	BH-01	30.0	gneiss	2.75	2750.0	64.35	25.75
4	BH-02	15.0	granite	2.63	2630.0	64.4	25.8
5	BH-02	25.0	schist	2.71	2710.0	64.4	25.8

innerjoin keeps only rows where the :borehole value exists in both tables. BH-03 has no samples, so it is dropped. If you want to keep all locations even when there are no matching samples, use leftjoin or outerjoin instead.

3.3.8 Handling missing data

Real data has gaps. Julia represents missing values with a special value called missing:

```
# Create data with a gap
obs = DataFrame(
    station = ["A", "B", "C", "D"],
    temp_C  = [12.3, missing, 14.1, 13.7]
)
obs
```

	station	temp_C
	String	Float64?
1	A	12.3
2	B	<i>missing</i>
3	C	14.1
4	D	13.7

```
# Drop rows with missing values
dropmissing(obs, :temp_C)
```

	station	temp_C
	String	Float64
1	A	12.3
2	C	14.1
3	D	13.7

```
# Compute mean, skipping missing values
using Statistics
mean(skipmissing(obs.temp_C))
```

```
13.366666666666665
```

`skipmissing` is a wrapper that tells functions to ignore the gaps. Without it, `mean` would return `missing` because any operation involving `missing` propagates the unknown.

3.4 Working with file paths

When your project has many data files, you need to build file paths correctly. Julia's `joinpath` function handles this across operating systems (so you don't have to worry about `/` vs `\`):

```
# Build a path
path = joinpath("data", "field_campaign", "gravity.txt")
println(path)
```

data/field_campaign/gravity.txt

Other useful path functions:

```
println("Current directory: ", pwd())
println("Files here: ", readdir("."))
```

```
# Check if a file exists before reading
if isfile("samples.csv")
    println("samples.csv exists!")
else
    println("File not found")
end
```

samples.csv exists!

Organise your data early

Create a data/ folder in your project from the start. Put raw data in data/raw/, processed data in data/processed/. This habit will save you from chaos later, especially when you have 50 stations and 3 field campaigns.

3.5 Searching, copying, and renaming files

In practice, geoscientific data rarely arrives in one tidy folder. You might receive a USB drive with hundreds of files scattered across nested directories, including raw measurements from field instruments, GPS logs, photos, metadata files, all mixed together. A common first task is: **find all files of a certain type, rename them consistently, and copy them into a single clean folder.**

Julia has everything you need for this built into the standard library, no extra packages required.

3.5.1 Listing files in a directory

```
# readdir returns the names of everything in a folder
for name in readdir(".")
    println(name)
end
```

```
01-getting-started-julia.html
01-getting-started-julia.qmd
01-getting-started-julia_files
02-files-data.html
02-files-data.qmd
02-files-data_files
03-plotting-maps.qmd
03-plotting-maps_files
04-neural-networks.qmd
gravity.txt
gravity_corrected.txt
samples.csv
stations.txt
```

By default `readdir` gives you the names without the full path. Pass `join = true` to get complete paths:

```
for path in readdir("data/raw", join = true)
    println(path)
end
```

3.5.2 Searching recursively with `walkdir`

`walkdir` is the key function for searching through directories and all their subdirectories. It yields a tuple `(root, dirs, files)` at each level:

```
# Find every .csv file anywhere inside data/
for (root, dirs, files) in walkdir("data")
    for f in files
        if endswith(f, ".csv")
            full_path = joinpath(root, f)
            println(full_path)
        end
    end
end
```

```

        end
    end
end

```

You can wrap this into a reusable function:

```

"""
    find_files(dir, ext)

    Recursively find all files with extension `ext` in `dir` and its subdirectories.
    Returns a vector of full paths.
"""
function find_files(dir, ext)
    results = String[]
    for (root, dirs, files) in walkdir(dir)
        for f in files
            if endswith(f, ext)
                push!(results, joinpath(root, f))
            end
        end
    end
    return results
end

```

Main.Notebook.find_files

```

# Example: find all .dat files under a field campaign folder
dat_files = find_files("data/raw", ".dat")
println("Found $(length(dat_files)) .dat files")

```

3.5.3 Extracting parts of a file path

Julia has several functions for splitting a path into its components:

```

example_path = joinpath("surveys", "2024", "gravity", "station_042.csv")

println("Directory: ", dirname(example_path))
println("Filename:  ", basename(example_path))

```

```
println("Name only: ", splitext(basename(example_path))[1])
println("Extension: ", splitext(basename(example_path))[2])
println("Split path: ", splitpath(example_path))
```

```
Directory: surveys/2024/gravity
Filename:  station_042.csv
Name only: station_042
Extension: .csv
Split path: ["surveys", "2024", "gravity", "station_042.csv"]
```

3.5.4 Copying and renaming files

Suppose you have collected EM data from multiple campaigns and the instrument saved files with unhelpful names like `meas_001.dat`. You want to copy them into a single processed/ folder, renaming each file to include the station name and date.

Here is a complete worked example. We first create a fake directory tree so you can run everything:

```
# --- Set up a fake directory tree ---
for station in ["EM01", "EM02", "EM03"]
    dir = joinpath("fake_survey", station)
    mkpath(dir)    # like mkdir -p: creates all intermediate directories
    for i in 1:3
        filepath = joinpath(dir, "meas_$(lpad(i, 3, '0')).dat")
        write(filepath, "fake data for $station measurement $i\n")
    end
end
println("Created fake survey tree:")
for (root, dirs, files) in walkdir("fake_survey")
    for f in files
        println(" ", joinpath(root, f))
    end
end
```

Created fake survey tree:

```
fake_survey/EM01/meas_001.dat
fake_survey/EM01/meas_002.dat
fake_survey/EM01/meas_003.dat
fake_survey/EM02/meas_001.dat
```

```
fake_survey/EM02/meas_002.dat
fake_survey/EM02/meas_003.dat
fake_survey/EM03/meas_001.dat
fake_survey/EM03/meas_002.dat
fake_survey/EM03/meas_003.dat
```

Now let's search for all .dat files, rename them to include the station name, and copy them into a clean output folder:

```
# --- Find, rename, and copy ---
output_dir = "collected_data"
mkpath(output_dir)

for (root, dirs, files) in walkdir("fake_survey")
  for f in files
    if endswith(f, ".dat")
      # The station name is the immediate parent folder
      station = basename(root)

      # Build a new name: EM01_meas_001.dat
      new_name = station * "_" * f

      src = joinpath(root, f)
      dst = joinpath(output_dir, new_name)

      cp(src, dst, force = true) # copy; force=true overwrites if exists
    end
  end
end

# Check the result
println("Files in $(output_dir):")
for f in sort(readdir(output_dir))
  println(" ", f)
end
```

Files in collected_data:

```
EM01_meas_001.dat
EM01_meas_002.dat
EM01_meas_003.dat
EM02_meas_001.dat
```



```

EM02_meas_002.dat
EM02_meas_003.dat
EM03_meas_001.dat
EM03_meas_002.dat
EM03_meas_003.dat

```

3.5.5 Key functions at a glance

Function	What it does
<code>readdir(dir)</code>	List contents of a directory
<code>readdir(dir, join=true)</code>	Same, but returns full paths
<code>walkdir(dir)</code>	Recursively walk through all subdirectories
<code>mkpath(dir)</code>	Create a directory (and parents), like <code>mkdir -p</code>
<code>cp(src, dst)</code>	Copy a file
<code>mv(src, dst)</code>	Move or rename a file
<code>rm(path)</code>	Delete a file
<code>rm(path, recursive=true)</code>	Delete a directory and everything inside it
<code>isfile(path)</code>	Check if a path points to an existing file
<code>isdir(path)</code>	Check if a path points to an existing directory
<code>joinpath(parts...)</code>	Build a path from pieces (cross-platform)
<code>basename(path)</code>	Filename from a path
<code>dirname(path)</code>	Directory from a path
<code>splitext(name)</code>	Split into (name, ".ext")
<code>splitpath(path)</code>	Split into a vector of path components

3.6 Geoscientific file formats

Geoscience uses several specialised binary formats to store large, multidimensional datasets efficiently. Julia has packages for all the common ones. Here is a brief overview. Detailed usage will be added in future updates of this chapter.

3.6.1 NetCDF

NetCDF (Network Common Data Form) is the standard format for climate, ocean, and atmospheric data. Use the `NCDatasets.jl` package:

```
using NCDatasets
```

```
# Read a NetCDF file
ds = NCDataset("temperature.nc")
temp = ds["temperature"][ :, :, :] # read the full 3D array
lat  = ds["latitude"][ :]
lon  = ds["longitude"][ :]
close(ds)

# Or use the do-block pattern (auto-closes)
NCDataset("temperature.nc") do ds
    temp = ds["temperature"][ :, :, :]
    println("Shape: ", size(temp))
end
```

```
# Write a NetCDF file
NCDataset("output.nc", "c") do ds
    defDim(ds, "x", 10)
    defDim(ds, "y", 20)
    v = defVar(ds, "elevation", Float64, ("x", "y"))
    v[ :, :] = rand(10, 20)
end
```

3.6.2 HDF5

HDF5 (Hierarchical Data Format) is widely used in geophysics, remote sensing, and seismology. Use `HDF5.jl`:

```
using HDF5

# Write
h5open("model.h5", "w") do f
    f["resistivity"] = rand(100, 100, 50)
    f["depth"] = collect(0.0:0.5:24.5)
end

# Read
h5open("model.h5", "r") do f
    rho  = read(f, "resistivity")
    depth = read(f, "depth")
    println("Model size: ", size(rho))
end
```

```
end
```

3.6.3 SEG-Y

SEG-Y is the seismic industry standard. Use `SegyIO.jl`:

```
using SegyIO

# Read a SEG-Y file
block = segy_read("seismic_line.segy")
traces = block.data           # matrix of traces
headers = block.traceheaders  # trace header information
```

3.6.4 Shapefiles and GeoJSON

For vector geospatial data (points, polygons, map boundaries), use `Shapefile.jl` or `GeoJSON.jl`:

```
using Shapefile

table = Shapefile.Table("geology_map.shp")
# Access as a DataFrame
using DataFrames
df = DataFrame(table)
```

 These examples are not executed

The geoscientific format examples above use `eval: false` style. They show the code patterns but don't run here because they require actual data files. When you have your own NetCDF, HDF5, or SEG-Y files, copy the pattern and replace the file name.

 Want more formats or deeper coverage?

This section will grow as the book develops. If there is a specific geoscience format or workflow you need (GRIB, GeoTIFF, ASDF, miniSEED, EDI, ModEM, or anything else) please [open an issue on GitHub](#) and describe your use case. Your request will help us prioritise what to add next.

3.7 Clean up

Let's remove the temporary files we created in this chapter:

```
for f in ["stations.txt", "gravity.txt", "gravity_corrected.txt", "samples.csv"]
  isfile(f) && rm(f)
end
println("Cleaned up.")
```

Cleaned up.

4 Plotting & Data Visualisation

At this point you know how to write Julia code, work with functions and data structures, and read data from files into arrays and DataFrames.

In this chapter we will learn how to **visualise** that data. We start from the very basics, a single line on a pair of axes, and build up to the kinds of plots geoscientists use every day: scatter plots with error bars, borehole logs, contour maps, heatmaps, multi-panel figures, and more. By the end you will be able to turn any dataset into a publication-ready figure and save it as PNG, PDF, or SVG.

4.1 The Makie ecosystem

Julia has several plotting libraries. In this book we use **Makie**, a modern, high-performance plotting ecosystem designed for scientific visualisation. Makie is split into **backends**. You write the same plotting code, but choose a backend depending on where you want the output:

Backend	Package	Best for
CairoMakie	CairoMakie.jl	Static figures (PNG, PDF, SVG) for papers, reports, this book
GLMakie	GLMakie.jl	Interactive, GPU-accelerated windows for rotating 3-D models, exploring large datasets on your screen
WGLMakie	WGLMakie.jl	Interactive plots in a web browser or Jupyter/Pluto notebook

! CairoMakie vs GLMakie: which one do I use?

Think of Makie as one plotting language with different “printers”:

- **CairoMakie** renders to a file (PNG, PDF, SVG). There is no window. You call `save("fig.pdf", fig)` and get a crisp vector graphic ready for a journal. It works everywhere, including headless servers and HPC clusters with no display. This is what we use throughout this book.
- **GLMakie** opens an **interactive window** on your screen, powered by your GPU. You can

pan, zoom, and rotate 3-D scenes with the mouse, perfect for exploring a velocity model, a point cloud, or a seismicity catalogue before you decide which view to export. You need a display and a working OpenGL driver.

- **WGLMakie** does the same interactive thing, but inside a **web browser**. It is the best choice for Jupyter or Pluto notebooks.

The key point: **your plotting code does not change**. You just swap using CairoMakie for using GLMakie (or using WGLMakie). So learn the API once, and pick the backend that fits the task.

For this book we use **CairoMakie** because it produces crisp vector graphics and does not need a GPU. When you start exploring 3-D geological models or large point clouds interactively, switch to **GLMakie**. The plotting code stays the same, only the using line changes.

💡 What about Plots.jl and Plotly?

Plots.jl is another popular Julia plotting package with a different API. It supports multiple backends (GR, Plotly, PGFPlotsX, etc.) through a single interface. If you prefer Plotly-style interactive HTML plots, you can use `Plots.jl` with the `plotlyjs()` backend, or use **PlotlyJS.jl** directly. We chose Makie for this book because its composability (building complex figures from simple pieces) and performance with large datasets make it a natural fit for geoscientific work. But there is no wrong choice — explore both and use what feels right for your workflow.

i GMT.jl: the Generic Mapping Tools for Julia

If you work with geographic or bathymetric data, you should know about **GMT.jl**, a Julia wrapper around the legendary **Generic Mapping Tools (GMT)**. GMT has been the gold standard for map-making in the earth sciences for over 30 years. The Julia wrapper gives you the full power of GMT (coastlines, topographic relief, geodetic projections, gravity grids, focal mechanisms) from within Julia.

GMT.jl is not covered in this chapter because it has its own extensive documentation and a different API from Makie. But if your work involves publication-quality maps with coastlines, plate boundaries, or bathymetry, it is absolutely worth exploring:

<https://www.generic-mapping-tools.org/GMT.jl/stable/>

4.2 Installing CairoMakie

Before we can plot anything, we need to add CairoMakie to our project environment. Press `]` in the REPL to enter Pkg mode:

```
pkg> add CairoMakie
```

Or from normal Julia code:

```
using Pkg; Pkg.add("CairoMakie")
```

You only need to do this once per environment. CairoMakie is already in this book's `Project.toml`, so if you cloned the book repository and ran `Pkg.instantiate()`, you are all set.

i First load is slow, that is normal

The very first time you run using CairoMakie in a fresh Julia session, it takes a while (sometimes 20 to 40 seconds) because Julia compiles the package. Subsequent calls in the same session are instant. This is Julia's "time to first plot" (TTFP). It has improved enormously over recent versions and continues to get faster.


```
using CairoMakie
CairoMakie.activate!(type = "png")
```

4.3 Your first plot

Let's start with the simplest possible plot, a line:

```
x = 0:0.1:10
y = sin.(x)

lines(x, y)
```




That is it. `lines(x, y)` creates a figure, adds an axis, and draws a line through the points. `Makeie` returns a `FigureAxisPlot` object that renders automatically.

4.3.1 Adding labels and a title

A plot without labels is just a pretty picture. Always label your axes:

```
fig, ax, plt = lines(x, sin.(x),
    axis = (
        xlabel = "Distance (km)",
        ylabel = "Amplitude",
        title = "Synthetic waveform"
    ),
    linewidth = 2,
    color = :steelblue
)
fig
```

The `fig, ax, plt = ...` pattern gives you handles to the **figure**, the **axis**, and the **plot object**. You will use these handles a lot when building more complex figures.

4.4 Scatter plots

Scatter plots are the workhorse of geoscience: plotting measurement locations, geochemical concentrations, geophysical anomalies, or any point data.


```
# Synthetic borehole collar locations
n = 30
easting = 500 .+ 200 .* rand(n)
northing = 7000 .+ 150 .* rand(n)
depth = 50 .+ 100 .* rand(n)

fig, ax, plt = scatter(easting, northing,
    color = depth,
    colormap = :viridis,
    markersize = 12,
    axis = (
```

```

    xlabel = "Easting (m)",
    ylabel = "Northing (m)",
    title = "Borehole collar locations",
    aspect = DataAspect()
)
)
Colorbar(fig[1, 2], plt, label = "Depth (m)")
fig

```



A few things to notice:


- **color = depth** maps a data vector to colour, just like you would colour-code sample values on a field map.
- **colormap = :viridis** picks a perceptually uniform colourmap. Other good choices for geo-science: **:turbo**, **:RdBu** (diverging, great for anomalies), **:terrain**, **:deep**.
- **Colorbar(...)** adds a colour legend. The position **fig[1, 2]** means “row 1, column 2 of the figure layout.”
- **aspect = DataAspect()** keeps the x and y scales equal — essential when plotting spatial coordinates.

4.5 Line plots with multiple series

Geoscientists often compare several datasets on the same axes: different field campaigns, model predictions versus observations, or time series from multiple stations.

```
time = 0:0.5:50          # time in seconds
signal_1 = sin(0.5 .* time) .+ 0.1 .* randn(length(time))
signal_2 = sin(0.5 .* time .+ pi/3) .+ 0.1 .* randn(length(time))

fig = Figure(size = (700, 400))
ax = Axis(fig[1, 1],
    xlabel = "Time (s)",
    ylabel = "Voltage (mV)",
    title = "EM receiver channels"
)
lines!(ax, time, signal_1, label = "Rx-1", linewidth = 2)
lines!(ax, time, signal_2, label = "Rx-2", linewidth = 2, linestyle = :dash)
axislegend(ax, position = :rt)
fig
```



Notice the ! in lines! and axislegend!-style calls. The exclamation mark means “add to an existing axis” rather than creating a new one. This is a consistent Makie convention.


4.6 Error bars

Field measurements always have uncertainty. Use `errorbars!` to show it:

```
stations = 1:8
gravity = [12.3, 14.1, 13.5, 15.8, 16.2, 14.9, 13.7, 12.8]
uncert = [0.5, 0.8, 0.4, 0.6, 0.9, 0.3, 0.7, 0.5]

fig = Figure(size = (600, 400))
ax = Axis(fig[1, 1],
    xlabel = "Station",
    ylabel = "Gravity anomaly (mGal)",
    title = "Profile A - Bouguer anomaly",
    xticks = 1:8
)
```

```
errorbars!(ax, stations, gravity, uncert, whiskerwidth = 8)
scatter!(ax, stations, gravity, markersize = 10, color = :black)
fig
```




4.7 Bar charts and histograms

4.7.1 Histogram of sample data

```
# Simulated grain sizes (log-normal distribution)
grain_sizes = exp.(randn(500) .* 0.6 .+ 1.0)

fig = Figure(size = (600, 400))
ax = Axis(fig[1, 1],
    xlabel = "Grain size (mm)",
    ylabel = "Count",
    title = "Grain-size distribution")
```


```
)
hist!(ax, grain_sizes, bins = 30, color = (:steelblue, 0.7))
fig
```



4.7.2 Bar chart

```
minerals = ["Quartz", "Feldspar", "Mica", "Olivine", "Pyroxene"]
abundance = [35, 25, 15, 15, 10]

fig = Figure(size = (600, 400))
ax = Axis(fig[1, 1],
    xlabel = "Mineral",
    ylabel = "Abundance (%)",
    title = "Modal composition",
    xticks = (1:5, minerals)
)
barplot!(ax, 1:5, abundance, color = :teal)
```

`fig`


4.8 Heatmaps and contour plots

Gridded data is everywhere in geoscience: gravity grids, magnetic surveys, temperature fields, elevation models. Makie handles these with `heatmap` and `contour/contourf`.

4.8.1 Heatmap


```
# Synthetic gravity anomaly on a 2-D grid
x_grid = range(0, 10, length = 60)
y_grid = range(0, 8, length = 50)
gz = [10 * exp(-((xi - 4)^2 + (yi - 3)^2) / 3) -
      5 * exp(-((xi - 7)^2 + (yi - 6)^2) / 2)
      for xi in x_grid, yi in y_grid]
```

```
fig = Figure(size = (700, 500))
ax = Axis(fig[1, 1],
          xlabel = "Easting (km)",
          ylabel = "Northing (km)",
          title = "Bouguer gravity anomaly",
          aspect = DataAspect()
        )
hm = heatmap!(ax, x_grid, y_grid, gz, colormap = :RdBu)
Colorbar(fig[1, 2], hm, label = " $\Delta g_B$  (mGal)")
fig
```



4.8.2 Filled contour plot

```
fig = Figure(size = (700, 500))
ax = Axis(fig[1, 1],
          xlabel = "Easting (km)",
          ylabel = "Northing (km)",
          title = "Gravity anomaly (contour)",
          aspect = DataAspect()
)
co = contourf!(ax, x_grid, y_grid, gz, levels = 15, colormap = :RdBu)
contour!(ax, x_grid, y_grid, gz, levels = 15, color = :black, linewidth = 0.5)
Colorbar(fig[1, 2], co, label = " $\Delta g_B$  (mGal)")
fig
```



Overlaying contour! lines on a contourf! fill is a classic geophysics style. The thin black lines make it easier to read precise values.

4.9 Multi-panel figures (subplots)

Scientific papers often require several panels in one figure. Makie uses a **grid layout** system: `fig[row, col]` places content anywhere on the grid.

```
fig = Figure(size = (700, 500))

# Panel a - line plot
ax1 = Axis(fig[1, 1], title = "(a) Signal", xlabel = "Time (s)", ylabel = "mV")
```


```
lines!(ax1, time, signal_1)

# Panel b – scatter
ax2 = Axis(fig[1, 2], title = "(b) Collars",
           xlabel = "Easting", ylabel = "Northing", aspect = DataAspect())
scatter!(ax2, easting, northing, color = depth, colormap = :viridis, markersize = 8)

# Panel c – histogram
ax3 = Axis(fig[2, 1], title = "(c) Grain sizes", xlabel = "mm", ylabel = "Count")
hist!(ax3, grain_sizes, bins = 25, color = (:teal, 0.7))

# Panel d – contour
ax4 = Axis(fig[2, 2], title = "(d) Gravity", xlabel = "Easting (km)",
           ylabel = "Northing (km)", aspect = DataAspect())
contourf!(ax4, x_grid, y_grid, gz, levels = 12, colormap = :RdBu)

fig
```



You can span rows or columns for insets and wide panels — see the [Makie layout tutorial](#) for advanced arrangements.


4.10 Customising appearance

4.10.1 Fonts and themes

Makie ships with several built-in themes. You can also override individual settings:

```
with_theme(theme_minimal()) do
  fig = Figure(size = (600, 400))
  ax = Axis(fig[1, 1],
```

```
    xlabel = "Frequency (Hz)",
    ylabel = "Power spectral density",
    title = "Minimal theme example",
    xscale = log10,
    yscale = log10
)
freqs = 10.^ range(-1, 3, length = 200)
psd = 1 ./ freqs .+ 0.01 .* abs.(randn(200))
lines!(ax, freqs, psd, linewidth = 2)
fig
end
```



Other themes to try: `theme_dark()`, `theme_light()`, `theme_black()`, `theme_ggplot2()`.

4.10.2 Colourmaps

Choosing the right colourmap matters, especially in geoscience where a bad colourmap can hide features or mislead the reader. Some guidelines:

Data type	Good choices	Avoid
Sequential (elevation, depth, concentration)	:viridis, :cividis, :deep, :thermal	:jet (perceptual jumps)
Diverging (anomalies centred on zero)	:RdBu, :BrBG, :PiYG	Any sequential map
Categorical (lithology, fault type)	:tab10, :Set1	Continuous maps
Terrain / bathymetry	:terrain, :topo	:hot, :gray


```
fig = Figure(size = (650, 280))

ax1 = Axis(fig[1, 1], title = ":viridis (sequential)", aspect = DataAspect())
heatmap!(ax1, gz, colormap = :viridis)

ax2 = Axis(fig[1, 2], title = ":RdBu (diverging)", aspect = DataAspect())
heatmap!(ax2, gz, colormap = :RdBu)

ax3 = Axis(fig[1, 3], title = ":terrain", aspect = DataAspect())
heatmap!(ax3, gz, colormap = :terrain)

fig
```




4.11 Log-scale axes

Geophysical data often spans many orders of magnitude: resistivity, frequency, spectral power. Use `xscale = log10` or `yscale = log10`:

```
freq = 10.^ range(-2, 5, length = 50)
apparent_res = 100 .* (1 .+ 0.3 .* sin.(log10.(freq) .* 2))

fig = Figure(size = (600, 400))
ax = Axis(fig[1, 1],
    xlabel = "Frequency (Hz)",
    ylabel = "Apparent resistivity ( $\Omega\cdot\text{m}$ )",
    title = "MT sounding curve",
    xscale = log10,
    yscale = log10
)
lines!(ax, freq, apparent_res, linewidth = 2, color = :navy)
fig
```



4.12 Annotations and markers


Labelling key features on a plot (a fault, an anomaly peak, a sample location) is something you will do often:

```
fig = Figure(size = (600, 400))
ax = Axis(fig[1, 1],
          xlabel = "Easting (km)", ylabel = "Gravity (mGal)",
          title = "Annotated profile"
)
profile_x = range(0, 10, length = 100)
profile_g = 8 .* exp.-((profile_x .- 4) ./ 1.5).^2) .- 2
lines!(ax, profile_x, profile_g, linewidth = 2)

# Mark the peak
scatter!(ax, [4.0], [6.0], markersize = 14, color = :red)
text!(ax, 4.3, 6.2, text = "Anomaly peak", fontsize = 14, color = :red)

# Mark a fault
vlines!(ax, [7.0], color = :grey40, linestyle = :dash, linewidth = 1.5)
text!(ax, 7.2, -1.5, text = "Fault F1", fontsize = 12, color = :grey40)

fig
```

4.13 Saving figures

Use `save` to export your figure. The file extension determines the format:

```
save("gravity_profile.png", fig)           # raster (300 dpi by default)
save("gravity_profile.pdf", fig)           # vector, best for papers
save("gravity_profile.svg", fig)           # vector, good for web
save("gravity_profile.png", fig, px_per_unit = 3) # high-res raster (900 dpi)
```

💡 PDF for papers, PNG for presentations

For journal submissions, **PDF** or **SVG** gives you lossless vector graphics that look sharp at any zoom. For slides and posters, **PNG at high resolution** (`px_per_unit = 3` or higher) is more portable.

4.14 Putting it together: a geoscience figure

Let's combine several techniques into a single multi-panel figure typical of a geophysical report:

```
using Statistics


fig = Figure(size = (700, 550))

# --- Panel a: profile with error bars ---
ax1 = Axis(fig[1, 1:2], title = "(a) Gravity profile",
            xlabel = "Distance (km)", ylabel = "Bouguer anomaly (mGal)")
dist = range(0, 12, length = 25)
obs = 8 .* exp.-((dist .- 5) ./ 2).^2) .+ 0.3 .* randn(25)
uncert = 0.2 .+ 0.1 .* rand(25)
errorbars!(ax1, collect(dist), obs, uncert, whiskerwidth = 6)
scatter!(ax1, collect(dist), obs, markersize = 8, color = :black, label = "Observed")
model_g = 8 .* exp.-((dist .- 5) ./ 2).^2)
lines!(ax1, collect(dist), model_g, linewidth = 2, color = :red, label = "Model")
axislegend(ax1, position = :rt)

# --- Panel b: contour map ---
ax2 = Axis(fig[2, 1], title = "(b) Anomaly map",
            xlabel = "Easting (km)", ylabel = "Northing (km)",
            aspect = DataAspect())
co = contourf!(ax2, x_grid, y_grid, gz, levels = 12, colormap = :RdBu)
contour!(ax2, x_grid, y_grid, gz, levels = 12, color = :black, linewidth = 0.4)
Colorbar(fig[2, 0], co, label = "mGal", flipaxis = false)

# --- Panel c: histogram of residuals ---
residuals = obs .- model_g
ax3 = Axis(fig[2, 2], title = "(c) Residuals",
            xlabel = "Residual (mGal)", ylabel = "Count")
hist!(ax3, residuals, bins = 10, color = (:steelblue, 0.7))
vlines!(ax3, [0.0], color = :red, linestyle = :dash)
vlines!(ax3, [mean(residuals)], color = :black, linewidth = 2, label = "mean")
axislegend(ax3, position = :lt)

fig
```



4.15 Geographic maps with GeoMakie

So far every plot has used plain numeric axes — “Easting (km)” and “Northing (km).” But geoscientists often need to plot data on a **map** with coastlines, country borders, and a proper geographic projection. That is what **GeoMakie.jl** gives you.

GeoMakie provides a **GeoAxis**, a drop-in replacement for Makie’s **Axis** that understands coordinate reference systems (CRS) and projections. You feed it longitude/latitude data and tell it which projection you want; it handles the transformation and draws graticules automatically.

4.15.1 Installing GeoMakie


```
pkg> add GeoMakie
```

Or: using `Pkg; Pkg.add("GeoMakie")`. GeoMakie pulls in Proj.jl (for map projections) and NaturalEarth (for coastlines and country boundaries).

```
using GeoMakie
```

4.15.2 A world map in three lines

```
fig = Figure()
ga = GeoAxis(fig[1, 1], dest = "+proj=natearth", title = "Natural Earth projection")
lines!(ga, GeoMakie.coastlines(), color = :black, linewidth = 0.5)
fig
```



GeoAxis accepts any **PROJ-string** through the `dest` keyword. The default source CRS is WGS84 (longitude/latitude). Try replacing `"+proj=natearth"` with `"+proj=ortho"` for an orthographic globe view, or `"+proj=moll"` for a Mollweide equal-area projection.

4.15.3 Overlaying polygons and raster data

GeoMakie integrates with **GeoInterface.jl**, so you can plot any GeoInterface-compatible geometry: shapefiles, GeoJSON, geological unit polygons, fault traces, concession boundaries:

```
using GeoMakie, GeoJSON

# Load a GeoJSON of geological units (hypothetical file)
geol = GeoJSON.read("geology_units.geojson")

fig = Figure()
ga = GeoAxis(fig[1, 1], dest = "+proj=utm +zone=18 +south")
poly!(ga, geol, color = :lightgrey, strokecolor = :black, strokewidth = 0.5)
fig
```

💡 Common projections for geoscience

Region / Use case	PROJ string	Notes
Global equal-area	<code>+proj=moll</code>	Mollweide, good for thematic world maps
Global for visuals	<code>+proj=natearth</code>	Natural Earth, pleasant for presentations
Globe view	<code>+proj=ortho +lon_0=0 +lat_0=30</code>	Orthographic, looks like a photo from space
Local field survey	<code>+proj=utm +zone=35 +datum=WGS84</code>	Metric grid, precise local coordinates

Continental	<code>+proj=lcc +lat_1=30 +lat_2=60 +lon_0=10</code>	Lambert Conformal Conic, preserves shape
Polar	<code>+proj=stere +lat_0=90 +lon_0=0</code>	Stereographic, Arc- tic/Antarctic maps

You can browse all available projections at <https://proj.org/operations/projections/>.

4.15.4 Learn more about GeoMakie

- Documentation & gallery: <https://geo.makie.org/>
- Source code: <https://github.com/MakieOrg/GeoMakie.jl>

4.16 Where to learn more

- Makie documentation: <https://docs.makie.org/> with tutorials, gallery, and the full API reference.
- Beautiful Makie: <https://beautiful.makie.org/>, a gallery of polished example figures you can copy and adapt.
- Makie colour maps: the `ColorSchemes.jl` package gives you access to hundreds of colour maps. Browse them at <https://juliagraphics.github.io/ColorSchemes.jl/stable/>.

5 Neural Networks

At this point you know how to handle data files, manipulate tables, and create plots.

In this chapter we will build **neural networks** that learn patterns directly from data. We start with a hands-on example (generate data, train, evaluate) before unpacking the theory. The framework is **Lux.jl**, a modern Julia library that keeps model parameters and state explicit.

5.1 Your first neural network: Sauna Satisfaction Predictor

Let's skip the theory for now and train a neural network end to end. Data goes in, a model learns, predictions come out. Everything fits on your laptop in under a minute.

5.1.1 The problem

In an alternate universe, the Geological Survey of Finland has an office sauna in every building, and geologists are allowed to use it after a single qualifying event: surviving a long meeting with geophysicists.

Since this is a government office sauna, a visitor log is maintained with the seriousness usually reserved for core samples. For each sauna visit the log records three things:

- Sauna temperature (°C)
- Outside temperature (°C)
- Minutes since last coffee

After the session the visitor assigns a **Satisfaction Score** between 0 and 1, where 0 means “this was a management idea” and 1 means “I love my job.”

We will train a small neural network to predict the Satisfaction Score from these three inputs. This is a **regression** problem (the output is a continuous number) rather than classification, so we will use mean squared error (MSE) as the loss function.

5.1.2 The training data

The logbook has 200 visits. The hidden satisfaction rule works like this: satisfaction peaks around 80 °C, gets a boost when it is freezing outside, drops when the last coffee was too long ago, and includes random noise because people are unpredictable. The relationships are non-linear, so a straight line cannot capture them, but a small neural network can.

5.1.3 Setting up packages

```
using Lux, Random, Optimisers, Zygote, Statistics, Printf
```

Installing Lux

If you have not installed Lux yet, open the Julia REPL and run:

```
using Pkg
Pkg.add(["Lux", "Optimisers"])
```

5.1.4 Generate the visitor log

The logbook has been transcribed into a tidy text file called `sauna_log.txt`. Each row is one visit, with columns separated by commas. Here are the first few lines:

```
# Peek at the file
for (i, line) in enumerate(eachline("sauna_log.txt"))
    i > 6 && break
    println(line)
end
```

```
sauna_temp,outside_temp,coffee_min,score
78.2,-6.8,76.7,1.0000
66.9,8.8,8.2,0.1861
86.4,-21.4,49.2,0.9703
72.9,-29.5,119.1,0.5900
86.8,-23.7,90.2,0.5756
```


5.1.5 Read the data and normalise

We read the file, parse the numbers, and normalise each feature to zero mean and unit variance so the network treats all inputs equally:

```
lines = readlines("sauna_log.txt")
header = split(lines[1], ",")
data = reduce(hcat, [parse.(Float32, split(l, ",")) for l in lines[2:end]])

# rows: sauna_temp, outside_temp, coffee_min, score
X_raw = data[1:3, :]
Y      = data[4:4, :]

# Normalise inputs
mu = mean(X_raw, dims = 2)
sig = std(X_raw, dims = 2)
X = (X_raw .- mu) ./ sig

@printf "Visits: %d   Features: %s   Output: %s\n" size(X, 2) string(size(X)) string(size(Y))
```

Visits: 200 Features: (3, 200) Output: (1, 200)

5.1.6 Split into train and test

```
rng = Xoshiro(123)
n_train = Int(round(0.8 * size(X, 2)))
idx = randperm(rng, size(X, 2))

X_train = X[:, idx[1:n_train]]
Y_train = Y[:, idx[1:n_train]]
X_test  = X[:, idx[n_train+1:end]]
Y_test  = Y[:, idx[n_train+1:end]]

println("Train: ", size(X_train, 2), " Test: ", size(X_test, 2))
```

Train: 160 Test: 40

5.1.7 Build the model

One hidden layer, 8 neurons, tanh activation, and a single linear output. No sigmoid at the end because this is regression, not classification.

```
model = Chain(
    Dense(3 => 8, tanh),
    Dense(8 => 1)
)
```

```
Chain(
    layer_1 = Dense(3 => 8, tanh),          # 32 parameters
    layer_2 = Dense(8 => 1),                # 9 parameters
)      # Total: 41 parameters,
      #      plus 0 states.
```

In Lux the model is just a blueprint. Parameters and state are initialised separately:

```
ps, st = Lux.setup(rng, model)
```

```
((layer_1 = (weight = Float32[-0.71197426 0.5440084 -0.9027876; -0.62075335 -0.5090574 1.4643831; ... ; 0.716
```

This separation means you can save parameters, swap them between models, or pass them to a differential equation solver without any fuss.

5.1.8 Define the loss function

Mean squared error (MSE): predictions far from the truth get penalised quadratically.

```
function mse_loss(model, ps, st, data)
    x, y_true = data
    y_pred, st_new = model(x, ps, st)
    loss = mean((y_pred .- y_true) .^ 2)
    return loss, st_new, ()
end
```

mse_loss (generic function with 1 method)

5.1.9 Train the network

Adam optimiser, 500 epochs:

```
opt = Adam(0.01f0)
tstate = Training.TrainState(model, ps, st, opt)

for epoch in 1:500
    _, loss, _, tstate = Training.single_train_step!(
        AutoZygote(), mse_loss, (X_train, Y_train), tstate
    )
    if epoch == 1 || epoch % 100 == 0
        @printf "Epoch %4d  MSE = %.6f\n" epoch loss
    end
end
```

```
Epoch   1  MSE = 0.298768
Epoch 100  MSE = 0.021929
Epoch 200  MSE = 0.008098
Epoch 300  MSE = 0.006308
Epoch 400  MSE = 0.005393
Epoch 500  MSE = 0.005107
```

5.1.10 Evaluate on test data

```
Y_pred_test, _ = model(X_test, tstate.parameters, tstate.states)
test_mse = mean((Y_pred_test .- Y_test) .^ 2)
test_mae = mean(abs.(Y_pred_test .- Y_test))

@printf "Test MSE:  %.6f\n" test_mse
@printf "Test MAE:  %.4f  (average error in satisfaction score)\n" test_mae
```

```
Test MSE:  0.006323
Test MAE:  0.0615  (average error in satisfaction score)
```

5.1.11 Predict new scenarios

Now for the important part. HR wants to know: what sauna settings should they recommend after the next geophysicist meeting?

```

scenarios = [
    (sauna=80, outside=-20, coffee=10,
     label="Perfect: 80°C, -20°C, fresh coffee"),
    (sauna=65, outside=5, coffee=100,
     label="Sad: lukewarm, mild day, stale coffee"),
    (sauna=95, outside=-25, coffee=30,
     label="Extreme: 95°C, deep winter, recent coffee"),
    (sauna=75, outside=0, coffee=60,
     label="Average: 75°C, 0°C, 1 h since coffee"),
]

function predict_score(s, model, ps, st, mu, sig)
    x = Float32.([(s.sauna - mu[1]) / sig[1],
                  (s.outside - mu[2]) / sig[2],
                  (s.coffee - mu[3]) / sig[3]])
    pred, _ = model(reshape(x, 3, 1), ps, st)
    return pred[1]
end

println("Sauna Satisfaction Predictions")
println("-"^40)
for s in scenarios
    sc = predict_score(s, model,
                       tstate.parameters, tstate.states, mu, sig)
    @printf " %-38s %.2f\n" s.label sc
end
println("-"^40)

```

Sauna Satisfaction Predictions

Perfect: 80°C, -20°C, fresh coffee	1.04
Sad: lukewarm, mild day, stale coffee	0.07
Extreme: 95°C, deep winter, recent coffee	0.41
Average: 75°C, 0°C, 1 h since coffee	0.77

The predictions match intuition: 80 °C with freezing weather and fresh coffee scores high; a lukewarm room on a mild day with stale coffee scores low. HR can now issue sauna recommendations with full scientific confidence.

 Try it yourself

1. Add a fourth input, “meeting duration (minutes).” Does the network improve?
2. Make it deeper: $3 \Rightarrow 16 \Rightarrow 8 \Rightarrow 1$. Does it overfit on 200 samples?
3. Swap tanh for relu. How does the test error change?

5.2 What is a neural network?

A **neuron** multiplies inputs by weights, adds a bias, and passes the result through a non-linear activation:

$$z = \sigma(\mathbf{w}^\top \mathbf{x} + b)$$

Stack neurons side by side to get a **layer**; chain layers to get a **network**. Our `Chain(Dense(3  $\Rightarrow$  8, tanh), Dense(8  $\Rightarrow$  1))` is two layers.

5.2.1 How does it learn?

Four steps, repeated many times:

1. **Forward pass**: run data through the network.
2. **Loss**: measure how wrong the predictions are.
3. **Backward pass**: compute gradients of the loss with respect to every weight using automatic differentiation.
4. **Update**: nudge each parameter to reduce the loss (Adam, SGD, etc.).

That is the entire recipe.

 Work in progress

Some sections are still being written. The example below is complete and ready to run. Now that you have all the things in place, we will start with detailed discussions about machine learning especially scientific machine learning and learn how to use best of both: domain knowledge and data and machine learning to solve geoscientific problems.